

TASK 3-MiniKube Deployment Task

Step 1: Start the minikube server

Starting the minikube using “minikube start” command

```
vishal@LAPTOP-U45BV051:~$ minikube start --driver=docker --container-runtime=docker
minikube v1.35.0 on Ubuntu 24.04 (amd64)
Using the docker driver based on user configuration

The requested memory allocation of 2200MiB does not leave room for system overhead (total system memory: 2901MiB). You may face stability issues.
Suggestion: Start minikube with less memory allocated: 'minikube start --memory=2200mb'

Using Docker driver with root privileges
Starting "minikube" primary control-plane node in "minikube" cluster
Pulling base image v0.0.46 ...
Downloading Kubernetes v1.32.0 preload ...
> preloaded-images-k8s-v18-v1...: 333.57 MiB / 333.57 MiB 100.00% 1.24 Mi
> gcr.io/k8s-minikube/kicbase...: 500.31 MiB / 500.31 MiB 100.00% 1.61 Mi
Creating docker container (CPUs=2, Memory=2200MB) ...
Preparing Kubernetes v1.32.0 on Docker 27.4.1 ...
  • Generating certificates and keys ...
  • Booting up control plane ...
  • Configuring RBAC rules
  • Configuring bridge CNI (Container Networking Interface) ...
Verifying Kubernetes components...
  • Using image gcr.io/k8s-minikube/storage-provisioner:v5
Enabled addons: storage-provisioner, default-storageclass
kubectl not found. If you need it, try: 'minikube kubectl -- get pods -A'
Done! kubectl is now configured to use "minikube" cluster and "default" namespace by default
```

Step 2: Install kubectl

Install kubectl from the documentation code

```
vishal@LAPTOP-U45BV051:~$ curl -LO "https://dl.k8s.io/release/$(curl -L -s https://dl.k8s.io/release/stable.txt)/bin/linux/amd64/kubectl"
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload  Total   Spent    Left   Speed
100 138    100 138    0    0    219    0 --:--:-- --:--:-- --:--:--    219
100 54.6M 100 54.6M    0    0 1481k    0 0:00:37 0:00:37 --:--:-- 1895k
```

Step 3: Create deployment

Now create a deployment named r2 using the image ‘vishal15276/test1’

```
vishal@LAPTOP-U45BV051:~$ kubectl create deployment r2 --image=vishal15276t/test1 --port=80
deployment.apps/r2 created
```

Step 4: Verify the pods

Now give kubectl get pods to check if the container is running and wait until it starts running

```
vishal@LAPTOP-U45BV051:~$ kubectl get pods
NAME                                READY   STATUS              RESTARTS   AGE
r1-7b886b659-f2sv6                 1/1     Running             0          9m38s
r2-f784c9f59-7f7g9                 0/1     ContainerCreating   0          6s
```

```
vishal@LAPTOP-U45BV051:~$ kubectl get pods
NAME                                READY   STATUS              RESTARTS   AGE
r1-7b886b659-f2sv6                 1/1     Running             0          10m
r2-f784c9f59-7f7g9                 1/1     Running             0          46s
```

Step 5: Expose the deployment

Now expose the deployment using the expose command

```
vishal@LAPTOP-U45BV051:~$ kubectl expose deployment r2 --port=80 --type=NodePort
service/r2 exposed
vishal@LAPTOP-U45BV051:~$ minikube service r2
```

Step 6: Accessing the website

Now give service command to check the ip address of the deployed image

[illegible]

Step 7: Output page

The output will be displayed as follows

